BUILDER, STEP 20

DeepSeek V4.1 Flash read our ask two ways.

With a sharper ask, all 8 runs matched our count.

THE FIRST ASK

One line, more than one reading.

STEP 11, THE TASK

"Count every client and every route, and rank by the size of the change, up or down."

```
# step 11 pick: clients and routes pooled
route BE-BE: 25 -> 60 (+35)

# Opus 5.5: client and route pairs
1. Alpha Retail (example), BE-NL: 33 -> 0 (-33)
```

On step 11, Opus 5.5 asked which one we meant and recommended client and route pairs. These runs used print mode in the terminal, with nobody to answer. We did not test if that mattered.

THE FIRST ASK, THE RESULT

2 of 9 matched our count.

	MATCHED	COST PER RUN
Pi, DeepSeek V4.1 Flash	2 of 3	under 1 cent
Pi, Qwen3.8 Flash	0 of 2	under 2 cents
Claude Code, Opus 5.5	0 of 2	35 and 36 cents, as shown
Opus plans, DeepSeek builds	0 of 2	31 cents, as shown

The other 7 counted client and route pairs, a fair reading. The step 19 test went red on all 9. One lost run and one timed-out run are not counted.

THE SHARPER ASK

We added two sentences.

FIRST

"Count each client on its own and each route on its own, not client and route pairs."

SECOND

"Print one header line, then the three changes, each line in exactly this format:

```
1. route  AA-BB    10 ->   20  +10"
```

The same task, models and starting commit as before.

THE SHARPER ASK, THE RESULT

8 of 8 matched our count.

	MATCHED	STEP 19 TEST
Pi, DeepSeek V4.1 Flash	2 of 2	2 of 2
Pi, Qwen3.8 Flash	2 of 2	1 of 2
Claude Code, Opus 5.5	2 of 2	2 of 2
Opus plans, DeepSeek builds	2 of 2	2 of 2

The one red test had the right numbers and different spacing. In the plan row, the plan already held the three lines.

WHY NO PI STEPS

Pi has no built-in permission system.

→ Pi's README says it has no built-in permission system.

→ It runs with your permissions. It can run any command that you can run.

→ A new folder does not stop it. We ran it on our Mac in a partial sandbox, with made-up data. The network and most reads stayed open.

→ So this step gives no Pi steps. If you try it, use a spare computer or a container, made-up data only, and a key with a low credit limit. The results are the point.

WHAT THIS DOES NOT SHOW

A few runs. One small task.

- → It does not rank these models. Each model ran in one tool only.
- → Pi ran in a sandbox. Claude Code ran on our machine, with no sandbox. Per the logs, neither read our answer files. Our setup may have played a part.
- → We wrote the sharper ask after we knew the answer, and its format line says route.
- → Cost as shown is the figure Claude Code reports, plus Pi's cost in the plan row. We ran it signed in with a plan, so it is not a bill we paid.

FREE FILES

Get the files. Free.

- → The sharper ask
- → The results per setup
- → The steps to try it in Claude Code, with the test as your check. Mac steps, on a paid Claude plan. Tested 31 Oct 2026.

On my profile, under Featured: AI Starter Kit, step 20. Save this post.

Next: the Expert level starts with guardrails.